LESSON 6.2.3

Putting Principles into Practice

How to operationalize AI governance from decision rights and documentation to testing, monitoring, and incident response.

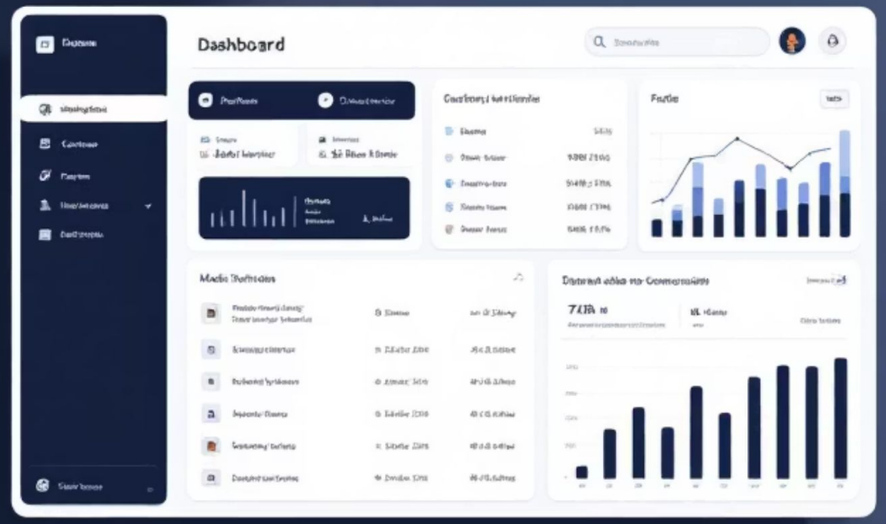

From Principles to Practice

Governance isn't theoretical — it requires concrete, repeatable operational steps.

01

Define Decision Rights

Who approves, monitors, and can pause a system?

02

Document the System

Model cards capture purpose, data, and limitations.

03

Test Before Release

Accuracy, fairness, stress tests, red-teaming.

04

Monitor After Deployment

Track technical and business metrics continuously.

05

Prepare for Incidents

Clear protocols for triage, communication, and correction.

Define Decision Rights

🟢 Approve

Who can authorize new AI deployments?

👁️ Monitor

Who tracks performance and security?

🔴 Pause / Roll Back

Who has authority to stop a system when issues arise?

The Model Card

A concise system profile — your AI system's source of truth.

- Purpose & Intended Users
- Data Sources & Training Period
- Known Limitations
- Prohibited Uses

Test Before You Deploy

Comprehensive pre-release testing catches problems before they reach users.

1

Accuracy Tests

Standard performance benchmarks

2

Fairness Checks

Evaluate across relevant demographic groups

3

Stress Tests

Unusual and edge-case inputs

4

Red-Teaming

Adversarial safety and security probing

Monitor After Release

⚙️ Technical Metrics

Error rates, model drift, latency, failure patterns

📊 Business Metrics

Complaint volumes, appeal rates, correction frequency

Monitoring isn't a one-time task — it's an ongoing commitment that spans the full lifecycle of the system.

Prepare for Incidents

Treat AI failures like any operational incident with clear, rehearsed protocols.

- Triage
- Communicate
- Patch or Retrain
- Tighten Guardrails

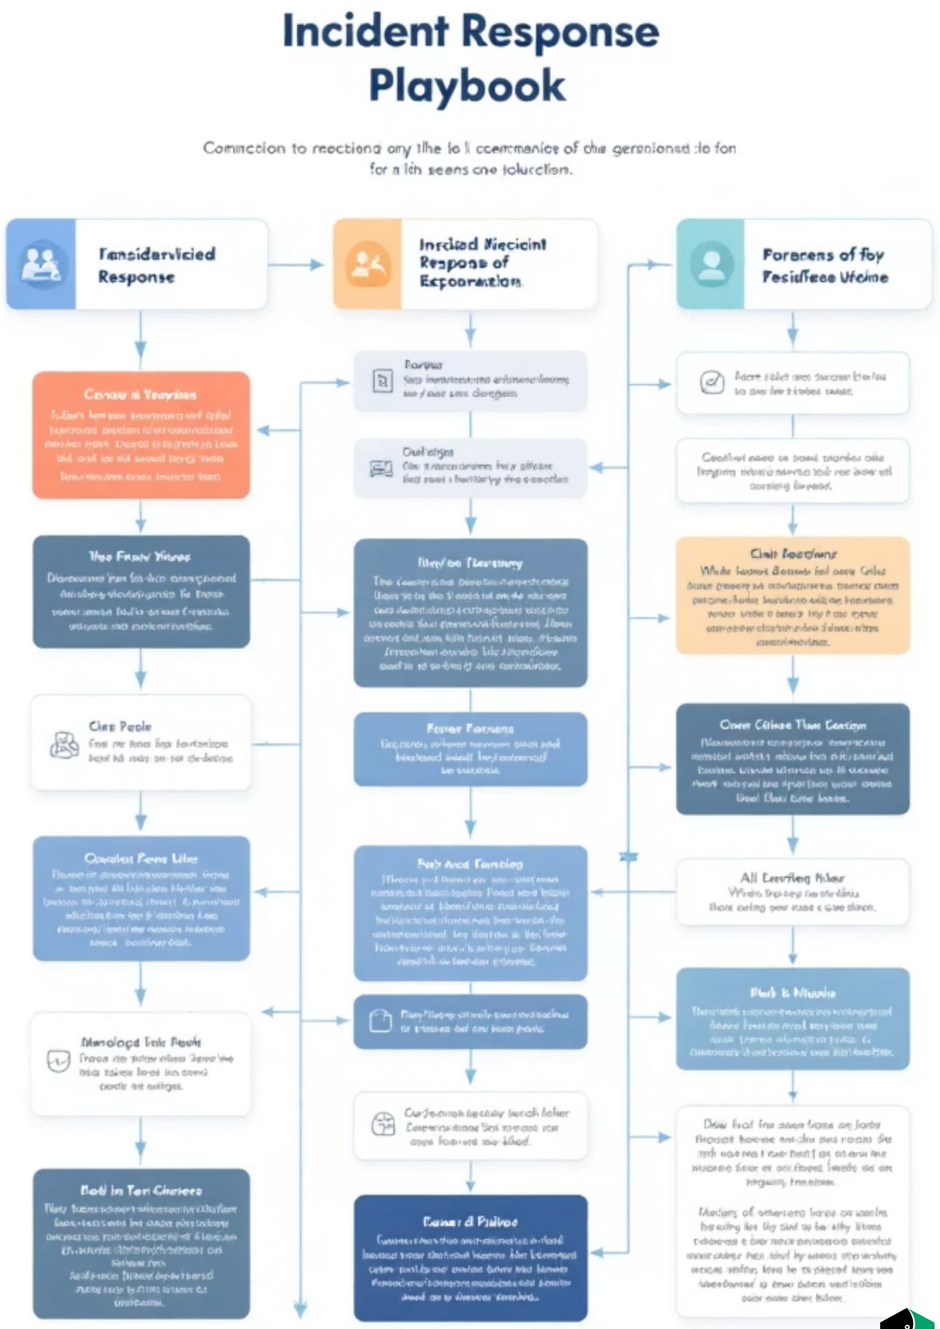

CASE STUDY

AI Writing Assistant: Governance in Action

An internal tool for drafting customer support emails, governed end-to-end.

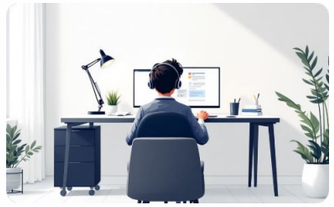

Ownership & Oversight

Customer support director owns it.

Security team reviews all updates.

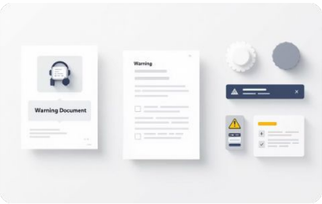

System Profile

Training sources documented. Legal notices flagged as prohibited use.

Pre-Release Testing

Verified: no fictitious discounts, no sensitive personal data in outputs.

Deployment & Monitoring: The Writing Assistant

🚦 Human-in-the-Loop Mandate

Agents must review all drafts until performance thresholds are met.

📋 Incident Playbook

Clear process to disable the assistant if unexpected outputs appear.

What Gets Monitored

Correction frequency

How often agents edit AI-generated drafts

Unusual output spikes

Flagged automatically for human review

Key Takeaways

Assign Decision Rights

Every AI system needs named owners for approval, monitoring, and rollback.

Document with Model Cards

Purpose, data, limitations, and prohibited uses in writing, before launch.

Test Thoroughly

Accuracy, fairness, stress, and adversarial red-teaming are non-negotiable.

Monitor Continuously

Track both technical drift and real-world business impact after deployment.

Have an Incident Plan

Triage, communicate, patch, retrain — know the playbook before you need it.